

Claude Code Bootstrap

Paste-into-Claude-Code specification to rebuild the Open Book research-desk playground from an empty folder. Everything below is an instruction to the building agent. Give Claude Code this document plus an empty git repo and say: "Build this."

Mission statement for the building agent. You are building **Open Book**: a static, no-build-step advisory-desk simulator with a two-sided shell — an **Advisor / Solutions view switch** over two tabs, **Idea Feed** (the daily sweep as a social-style origination feed with a Solutions curation layer) and **Advisor Book** (a client-tile grid embedding per-client portfolio reports) — in a J.P. Morgan Private Bank aesthetic, whose defining property is **transparent, derived, validated research discipline**: a goal-inference engine, a single idea→client fit engine, two conviction rubrics computed by a Python generator with hard validation, honest-dating via a fingerprint ledger, a tactical trigger gate, and sourced-or-stripped charts/technical. Build files in the order of §2. Do not invent alternative architectures; where this spec gives a constant, use exactly that constant. Where a detail is unspecified, choose the simplest option consistent with §1.

1 • Non-negotiable principles

- Idea→client links are **computed at render time** against the live client books — never hand-authored.
- Every score decomposes into visible named components (axes with notes, pillars with one-line reads, goal drivers in plain English).
- Conviction ("how good is the idea") and client-fit ("how right for this client") are separate engines, surfaced separately.
- Facts are tagged `sourced` (≥ 2 independent sources) or `estimated`; validators warn, never silently fix.
- Citation integrity: a named-author view citation requires a `basis` stating what was read and when; otherwise `ownView: true`. The generator rejects violations.
- Dates are honest: `postedAt` moves only when an idea's argument-fingerprint changes (committed ledger).
- Charts must carry fetched market data with provenance or they are stripped. Technicals come from a live scanner or stay tagged estimated.
- Derivatives-first expression: every idea's `preferredExpression` is an options/structured/OTC construction; direct lines are alternates. MiFID gating decides who can trade what.

2 • File manifest (build in this order)

#	File	Contents
1	<code>openbook.css</code> + <code>styles.css</code>	Design systems: the shell (light/dark, Source Serif 4 + Archivo) and the legacy pages (§9)
2	<code>data.js</code>	<code>window.SEED</code> : themes, seed ideas, 9 clients, synthetic sector history, complexity taxonomy (§3)
3	<code>goals.js</code>	<code>window.GOALS</code> : <code>deriveGoals</code> + helpers (§4)
4	<code>mapping.js</code>	<code>window.MAPPING</code> : <code>scoreIdeaForClient</code> + <code>flagClients</code> + tradability (§5)
5	<code>scanner.js</code>	11-rule portfolio scan + thin facades over MAPPING
6	<code>expressions.js</code>	<code>window.EXPRESSIONS</code> : structure knowledge base, keyword resolver (§7)

7	<code>build_today_focus.py</code>	Generator/validator: rubrics, gates, ledger (§6)
8	<code>tv_technicals.py</code> , <code>fetch_chart_series.py</code>	Enrichment scripts (§6.5)
9	<code>today_focus.json</code>	First sweep payload (author via §8 procedure)
10	<code>email.js</code> , <code>charts.js</code> , <code>methodology.js</code> , <code>morgan.js</code>	Support engines
11	<code>index.html</code> + <code>openbook.js</code>	THE SHELL (§7b) — header, view switch, theme, 𐌆K palette, Idea Feed, Solutions curation, Advisor Book grid + embeds
12	<code>classic.html</code> + <code>app.js</code> , <code>portfolio.html</code> , <code>onepager.html</code>	Legacy full app kept working in embed mode (masthead/tabs hidden via <code>?embed=1</code> , theme-synced); per-client portfolio report (what a book tile embeds); printable one-pager

3 • data.js contract

3.1 Themes (9)

ai (AI & Productivity) • power (Power & Infrastructure) • duration (Core Fixed Income)
broaden (Broadening Equity) • realassets (Real Assets) • resilience (Resilience & Protection)
gold (Gold & Currency) • health (Healthcare Innovation) • structured (Structured Outcomes)

3.2 Seed idea object

```
{ id, themeId, intent: add|protect|trim|income, title,  
  type: Thematic|Opportunistic|Strategic, assetClass, sector,  
  bucket: Growth|Income|Preservation, conviction: "High"|"Medium-High"|"Medium",  
  horizon, thesis, structures: [lead expression first – derivatives-first],  
  riskProfile?: {vol, beta, structured}, tradeStatement?, tickers? }
```

3.3 Client object (9 clients)

```
{ id, name, ccy, aum ($m), classification: Retail|Professional, mifid, relationship,  
  risk (free text – parsed), profile, split {assetClass: pct}, goals { objective,  
  horizon ("Long-term • 7-10 yrs" – parsed), funding? { headline ("Grow to $75m by 2034" /  
  "Fund $3.6m/yr..."), metricLabel, current, target, unit, status } },  
  liabilities?: [{name, amount, unit, note, rate?, near?}], positions: [{ name, ticker,  
  assetClass, sector, ccy, weightPct, pnlPct, entryDate, entrySpot, mat, note,  
  altKind?: uncorrelated|directional }], summary }  
INVARIANTS:  $\Sigma$  weightPct = 100 exactly; split  $\equiv \Sigma$  positions per assetClass.
```

3.4 The nine books (engineered test cases — reproduce these shapes)

id	Shape	Tests
aurora	Retail EUR €50m; MU 25.8% @ +295pnl; 72% USD; bond -15; gold+energy hedges; 10% cash	Concentration, FX mismatch, tax-harvest, Retail-OTC suppression
fable	Pro USD \$38m; NVDA 24/AVGO 14/MU 13/ASML 19/MSFT 18; 6% autocall; 4% cash	Options-fluent monetisation (collars/overwrites)
scott	Retail \$124m; 48% FI (one -24% long bond), dividend equity, PFE -38, O -11; 10% cash	Income mandate, bond swap, loss harvest, cash ladder
amar	Pro \$71m; BTC 16% @ -35; gold 12 +118; infra 19; EM 12; macro HF 7	Preservation-bucket-full (gold gap-fit = o), crypto rehab
jacob	Pro \$96m; balanced; liabilities \$19m mortgage + \$2.4m Q3 tax; 5% cash; UNH -26	Liability-driven Income/Preservation push
ben	Retail \$29m; 82% equity value/energy book (XOM 13, SHEL 10, BRK 18, AAPL 14, JPM 12, CAT 8, DGE 7 @ -53); only 6% FI, 4% real assets; income run-rate behind goal (\$0.5m vs \$0.8m/yr)	Income-gap + protection-gap rules, loss harvest (Diageo), non-complex-only expressions for Retail
prahnav	Retail \$44m; NVDA 22 +279, CEG 15 +181; buckets on target	Goals-aligned top-down, Retail-collar problem bottom-up
morgan ("Mitch")	Pro \$38m; NVDA 24, BTC 18, TSLA 11, 2x QQQ 10, venture 10; NO funding goal	Goal inference from stated-risk (aggressive) + book
tejpaul ("Fotis")	Retail \$60m; balanced 57/35; no risk profile, no goals, no liabilities	Pure revealed inference → moderate

3.5 Synthetic sector history

```
client.sectorHistory[sector] = 24 monthly numbers, index 0 = current allocation EXACTLY.  
Shape from dominant position pnl:  $\geq 50 \rightarrow$  'ramp' (start 12% of cur; index1 =  $\text{cur} \times 1.05$ );  
 $\leq -10 \rightarrow$  'fade' (start 130%); else 'core' (start 92%). Deterministic sine wobble  $\pm 4\%$  on  $t \geq 4$ .  
Cash excluded. Label clearly as replaceable seed data.
```

3.6 Complexity taxonomy (order matters: structured checked BEFORE otc)

```
STRUCTURED_KEYWORDS = [structured note, structured re-entry, buffered, autocall, autocallable,  
    acm+, phoenix, reverse convertible, revcon, fcn, range note, bren, certificate,  
    capital-protected, capital protected, participation note, memory coupon, barrier note,  
    credit-linked, twin-win, shark-fin, booster, halo]  
OTC_KEYWORDS = [collar, risk reversal, seagull, accumulator, decumulator, prepaid,  
    variable forward, fx forward, forward / collar, dcd, dcs, dual currency, covered call,  
    overwrite, cash-secured, call spread, protective put, strangle, straddle, fx put,  
    fx call, fx option, put spread]  
complexityOf(s): structured  $\rightarrow$  "structured" . otc  $\rightarrow$  "otc" . else "non-complex"  
Retail may trade structured NOTES (packaged securities); Retail may NOT trade OTC.
```

4 • goals.js — deriveGoals(client)

```
PARAMS = { base:8, nearBillW:1.0, debtBallastW:0.3, concPivot:15, concW:1.0,
  drawdownBump:15, shortHorizonW:3, incomeNeedW:8, serviceW:8, debtMatchW:0.6,
  reqReturnW:5, longHorizonW:3, willingW:30, willingWNoGoal:45,
  cautionProtW:28, cautionIncW:18,
  revealRiskW:0.5, revealDefenseW:0.3, revealConcW:0.2,
  fullDefenseAt:0.55, fullConcAt:0.25, serviceRate:0.06 }
PARSERS: horizon → {phase: drawdown if /drawdown|decumulat|retire/, years = midpoint of digits,
  default 7}
funding headline → incomeNeed ($m/yr via /[\$€£]?([\d.]+)m?\yr/), terminalTarget (via /to \$Xm/),
  years (via /by 20xx/ - 2026), requiredReturn = (target/base)^(1/years) - 1
liabilities → nearBill (one-off: /tax|due|bridge|bill|payable|q[1-4]|settlement/),
  service = debtLike(/mortgage|loan|line|credit|leverage|margin|borrow/) × rate (default 6%)
WILLINGNESS: stated - aggressive .95 · growth .80 · moderate .50 · conservative/income .30 · else
  null
revealed = .5×riskShare(Equity+Alts) + .3×(1 - min(1, defensive/.55)) + .2×min(1, topName%/25)
blend = stated==null ? revealed : .5 stated + .5 revealed
RAW SCORES (each bucket = base + named terms):
Preservation = 8 + 1.0·nearBill% + 0.3·debtLoad% + 1.0·max(0, topName%-15)
  + 15·drawdown + 3·max(0,5-H) + 28·(1-willingness)
Income = 8 + 8·incomeNeed% + 8·service% + 0.6·debtLoad% + 15·drawdown + 18·(1-willingness)
Growth = 8 + 5·reqReturn + 3·max(0,H-6) + (reqReturn>0 ? 30 : 45)·willingness
vector = largest-remainder normalise to 100. source/confidence: stated-goal .9 · stated-risk .7 ·
  revealed .55. Expose components + plain-English drivers. goalOverride short-circuits.
FOLDING currentBuckets: Equity+directional Alts → Growth; FI/Credit/RealAssets → Income;
  Cash/Commodity(gold)/uncorrelated Alts → Preservation; Structured by keyword purpose.
altKind explicit wins; else keyword-sniff (/uncorrelated|market-neutral|macro|multi-strat|hedge
  fund|
  managed futures|trend|cta|relative value|absolute return|long\short/ → Preservation).
goalAlignment(idea, client): gap = goalVector[bucketOfIdea] - currentBuckets[bucketOfIdea];
  mult = clamp(1 + gap/200, 0.85, 1.15).
```

5 • mapping.js — scoreIdeaForClient(idea, client)

```
PARAMS.weights (exact fractions): holdings .25/.85 ≈ .294 · mandate .25/.85 ≈ .294 ·
  gap .20/.85 ≈ .235 · concSector .15/.85 ≈ .176 (Σ = 1.00)
PARAMS.affinity: lambda .94; comfortByBucket {Growth:25, Income:15, Preservation:12} (by the
  ASSET's bucket, per-sector overrides allowed); penaltyPerPp 10, penaltyCap 100; noHolding 50.
  applyMin 45 · flagMin 50 · flagMax 6 · tiers ≥66 Strong / ≥48 Good / else Marginal.
AXES (each 0-100, with a one-line note):
1 AFFINITY = max(0, recency-weighted percentile (λ=.94) of current sector alloc within the
  24-mo sectorHistory - max(0, cur - sectorPeg)×10 capped 100)
2 MANDATE&RISK = .6×RiskSuitability(matrix riskProfile{vol,beta,structured} × mandateClass)
  + .4×IntentFit(matrix: growth↔appreciation 90 · income↔yield 90 · preservation↔protection 92 ·
  off-goal 40-75). mandateClass parsed from client.risk.
3 GAP = max(0, (goalTarget[bucket] - currentBucket)/goalTarget × 100); 0 when bucket full;
  fallback to sector headroom vs peg when no goal target. goalTarget = GOALS.goalsFor(client).
4 CONC-SECTOR: raw = (1 - HHI(normalised in-sector weights))×100; fitContribution = 100 - raw
  (invertForFit=true, flippable); no in-sector holdings → 50.
COMBINE: bracketFit = round(Σ score×weight); alignedFit = bracketFit × goalAlignment.mult;
tradability(idea,client) ∈ {0,1}: 0 iff Retail AND isOtc(naturalExpression); fit = trad ×
  alignedFit;
suppressed ideas surface WITH the MiFID reason (never silently dropped).
applies = trad && alignedFit ≥ 45; flagClients: fit ≥ 50, top 6.
Derived fallbacks: goalTypeOf(idea) from bucket/structures; riskProfileOf from sector beta lists
```

(HIGH_BETA: Tech/Crypto/Materials/Energy/Industrials/Consumer; LOW: Utilities/Gold/Rates/Credit/Infra/RealEstate/FX); naturalExpression = structures[0] if absent.

6 · build_today_focus.py — generator/validator

6.1 Payload shape

```
today_focus.json = { asOf, generated, sweep { sources[], rule, brief {summary, watch[3],
full[4-6]}, note }, earnings: [idea...], exEarnings: [idea...], meta {technicalAsOf,
technicalFeed} }
Idea fields: id, direction (long|short|neutral), stance (earnings only: pre-position|post-print),
tradeStatement (REQUIRED one-sentence instrument+direction+view), intent, action,
preferredExpression, name, ticker, sector, assetClass, bucket, goalType, riskProfile,
naturalExpression, headline, thesis, themeId|null, offThemWhy|null, ownView?, structures[],
variant {street, us, gap}, changeMyMind, conviction {pillars[]}, earnings {reportDate, reportWhen,
pre: impliedMovePct+impliedSource+historicalAvgMovePct+historicalSource+impliedVsHist / post:
result+resultSource+reaction+reactionSource, watch}, macro {metric, detail, source, watch} (ex-
earnings),
facts [{text, tag: sourced|estimated}], sources [{name, kind: data|news|calendar|analysis|view,
basis (REQUIRED when kind=view)}], tactical?, trigger?, triggered?, levels {tenor, entry, target,
stop?}, chart {kind: spark|band, unit, band?, refs[], caption, symbol, range?}
```

6.2 Conviction rubrics (computed, never authored)

EARNINGS (/8): asymmetry, consensus, catalyst, positioning — each 0-2, dq tag each.
Bands: High = raw≥6 & no zero & all sourced · "High — data gap" = same with ≥1 estimated ·
Medium = raw 4-5 or any zero or consensus unverified · Low = raw<4 or 2+ zeros (excluded).
EX-EARNINGS (/15): asymmetry 0-3, catalyst 0-2, consensus 0-2, thesis 0-2 (driver-specificity
0/1 + variant-view 0/1), houseview 0-2, positioning 0-2, technical 0-2.
score = raw/15×100; High ≥75 · Medium ≥55 · Watch ≥0.
GLOBAL CAP: any pillar dq ∈ {estimated, unverified} → label capped at Medium (ex-earnings) /
"High — data gap" (earnings).

6.3 Validators (warn loudly, never silently pass)

- ≥2 sources per idea; every fact tagged; tradeStatement present; view-citations must carry basis.
- Stance-aware dates: pre-position report ≥ asOf; post-print report < asOf AND result present with resultSource="sourced".
- Coverage: FX sector present; Rates present (sector Rates or assetClass Fixed Income); both earnings stances present; ≥1 tactical idea (has levels) and ≥1 tactical FX.
- Tactical gate: levels ⇒ trigger text + boolean triggered required; only triggered ship; dropped ones reported "held back".
- Chart integrity: strip any chart lacking series (≥2 points) + seriesSource.
- Intent defaulting; pillar count/keys/max checks; unknown-pillar warnings.

6.4 Fingerprint ledger (today_focus.ledger.json)

```
fp = sha1(json of: tradeStatement, headline, thesis, changeMyMind, direction, intent,
structures, fact texts, pillar [key, score, authoredNote||note], levels, earnings.result)
— excludes chart series / indicator timestamps / computed totals.
unseen id → postedAt = updatedAt = asOf · fp unchanged → carry both · fp changed → keep postedAt,
bump updatedAt. Warn when updatedAt > 10 days old. Ledger applied AFTER the tactical gate.
```

6.5 Enrichment scripts (stdlib only)

```
tv_technicals.py: POST scanner.tradingview.com/global/scan, columns [close, currency, SMA50,
SMA100, SMA200, RSI, BB.upper, BB.lower, Recommend.MA]. Rewrite ONLY the Technical pillar
(direction-aware: trend confirms long / vetoes short; neutral ideas scored on tape-calm), flip its
```

dq to sourced; add RSI colour to Positioning WITHOUT flipping its dq; rates yields skipped honestly; failures leave the authored pillar untouched. Macro symbol map: DXY→TVC:DXY, XAU→OANDA:XAUUSD, FX pairs→FX_IDC/OANDA, equities→NASDAQ:/NYSE:.

fetch_chart_series.py: GET query1.finance.yahoo.com/v8/finance/chart/{sym}?range={1mo|3mo|6mo|1y}&interval=1d → write series (rounded closes), seriesAsOf, seriesSource ("Yahoo Finance daily closes · SYM · range to DATE").

PIPELINE: author payload → tv_technicals.py → fetch_chart_series.py → build_today_focus.py → commit+push.

7 • expressions.js, email.js, scanner.js, methodology.js

- **expressions.js:** a dictionary of ~40 structures (zero-cost collar, call spread, put spread, risk reversal, strangle, covered call/overwrite, cash-secured puts, phoenix autocall, ACM+/HALO basket, reverse convertible, buffered note, capital-protected note, participation note, range accrual (BREN), dual-currency deposit, FX forward/collar, accumulator, bond ladder, gold accumulator, REIT/infra/utility baskets, liquid alternatives...). Each: label, complex flag, what/mechanics/example (realistic terms), pros[], cons[], goal-role scores {g,i,p}, context(client) personalisation. A keyword resolver maps free-text structure names → entries (substring matching, aliases). Every structures[] string used anywhere must resolve.
- **email.js:** idea × client → personalised letter: greeting, real-holding hook (name a position and its P&L), the desk view in client language, a concrete implementation block (structure, tenor, strikes, size), a tax/cash/FX action if the scanner flags one, a balanced risk paragraph, disclosure footer. One code path shared by all render surfaces.
- **scanner.js:** 11 severity-ranked findings rules over a book: single-name concentration >15%, underwater tax-harvest candidates, idle cash >8%, FX mismatch >30% non-base, missing protection with <5% Preservation vs goal, long-duration loss, liability-vs-cash squeeze, crypto drawdown, overwrite candidates (+50% winners), thin fixed income, missing income vs goal. Plus facades ideaFit/clientsForIdea delegating to MAPPING.
- **methodology.js:** builds `window.METHODOLOGY` by READING the live PARAMS objects (never copy constants) so the in-app "how scoring works" and the Morgan assistant cannot drift from code.

7b • The shell (index.html + openbook.js) — UX contract

- **Header:** "Open Book · PRIVATE BANK" lockup; an Advisor/Solutions **view switch**; a light/dark theme toggle (localStorage key `ob-theme`, applied before first paint, synced into the embedded iframe and across pages via storage events); a "Search everything 🔍" command palette over ideas + clients; an as-of stamp; two tabs: **Idea Feed**, **Advisor Book**.
- **Idea Feed layout:** Market Brief card (sweep summary, expandable to the watch bullets + full narrative) above a three-column grid — left rail: **TOP 3** tiles ranked by BREADTH (count of clients the live engine flags) then conviction, with a sweep-pinned idea leading as DESK PICK (label only — scores untouched), tap-to-focus; below it the **CLIENT TOOLKIT** (tap a client → their email builder). Middle: search input, asset-class chips, client-filter dropdown, count line, the feed. Right rail: **Ask your book** — the Morgan AI panel answering only from the live engines ("it never invents a number").
- **Idea card anatomy:** author identity from provenance (themeId set → "Solutions Desk"; themeId null → "Claude's Views"; advisor draft → "You — Solutions view" + DRAFT pill, no conviction ring); "posted Nd ago" reads the idea's OWN ledger postedAt (never the sweep date) with an "updated" chip when updatedAt differs; the CONV ring = conviction.score / 100 (click → scoring modal); ticker/asset tags; name; thesis; RECOMMENDATION block = tradeStatement, overridable by a Solutions implementation edit; inline chart (charts.js) only when a sourced series shipped; "Via ..." line listing ONLY kind="view" sources, else an "Own view — no outside research leaned on" pill when ownView; ♥ save with engagement counters (liked/drafted/checked by N advisors, localStorage-bumped); actions: **Idea Suitability** (per-client flowchart from scoreIdeaForClient — flagged clients with fit/tier + plain-English why, suppressed clients WITH their MiFID reason) and **Email** (email.js composer, drafts downloadable).
- **Solutions view (curation):** sees the whole board; Advisor view sees only APPROVED ideas. Curation bar: " + Add idea" composer (name, ticker, asset class, goal bucket, direction, thesis, one-sentence recommendation, comma-separated expressions → a custom idea mapped by the same fit engine, never desk-scored), restore-removed. Per-card: ✓ Approve / ✗ Discard toggle (custom ideas locked approved), ✎ Edit implementation (persisted per idea; the edited text feeds cards and emails), ✕ remove. All curation state in localStorage (`ob_ui_v1` : approval, implEdits, drafted/checked counters, viewMode; user store: customIdeas, dismissedFocus). The Advisor view renders ZERO curation UI.
- **Advisor Book tab:** a client-tile grid (name, headline stats); tap → hide the grid, show an iframe at `portfolio.html?embed=1&client=<id>` with an "← All clients" bar. Legacy pages detect `embed=1`, hide their own masthead/tab bar, add a back pill, and inherit the theme.

8 • The daily sweep procedure (operating loop)

1. Verify market state across five workstreams with ≥ 2 independent sources per fact: (a) Fed/rates/macro, (b) energy/gold/commodities, (c) FX/central banks, (d) US equities/earnings + a DISLOCATION SCAN (any widely-held large cap $>15\%$ off highs is a candidate), (e) global markets + the named-author list.
2. Named authors — core every sweep: Citrini, Serenity, TSCS, Brent Donnelly, Cem Karsan, Joseph Wang, SemiAnalysis. Conditional by domain: The Transcript (earnings), Benn Eifert/Moontower/Tier1 Alpha (vol), Fabricated Knowledge (semis), Rory Johnston/Doomberg (commodities), HighYield Harry (credit), The Market Ear/Brad Setser (flows), Andy Constan/Politano (macro). Cite ONLY what was actually read, with basis; otherwise ownView. Restated theses $\neq$ re-citation.
3. Re-underwrite the board: retire played/stale ideas (delete from payload), re-ground keepers (same id — the ledger preserves postedAt), add new ideas. Both earnings stances; FX + rates present; ≥ 1 triggered tactical with levels; every idea derivatives-first.
4. Write the sweep block: summary, 3 watch bullets, 4-6 full paragraphs, and a desk note archiving WHAT CHANGED (retired/new/re-grounded/gate decisions/dislocations scanned/author discipline/attribution line/discrepancies logged).
5. Run the three scripts; fix every validator warning; commit (chore: market sweep refresh YYYY-MM-DD (one-line desk summary)) and **push** (a local commit is invisible — the site deploys from the remote).

9 • Design system

```
SHELL (openbook.css) — light + dark themes, [data-theme="dark"] on <html>:  
--ground #FBFAF7 • --surface #FFF • --ink #17150F • --charcoal #1E1B16 • --brown #54301A  
--bronze #996F3D • --bronze-l #C7A97A • --gold #EBC98D (dark-mode accent) • --green #2FBF71  
--green-pos #2C7A4B • --brick #9C4A3A • hairlines #E5E1DA / #F0ECE4 • creams #F4F1EA / #EEEEAE  
Type: 'Source Serif 4' (serif display) + 'Archivo' (UI sans), Google Fonts. Texture: uppercase  
letter-spaced kickers, colour strip-bars on cards, SVG conviction rings, soft card shadows.  
Theme: localStorage 'ob-theme', applied PRE-PAINT (inline head script), synced to the iframe  
and across pages via storage events.  
LEGACY pages (styles.css — classic/portfolio/onepager): espresso #29211A on warm paper,  
bronze #9A7B4F, hairline rules, printed-research-note feel, with its own dark mapping.  
Goal colours: Growth dark • Income bronze • Preservation green.
```

10 • Acceptance tests (run before calling it done)

1. Serve the folder; both tabs render in both views and both themes; no console errors.
2. Gold idea $\times$ Amar $\rightarrow$ Gap fit o (Preservation full); $\times$ Scott $\rightarrow$ high Gap fit.
3. A zero-cost-collar idea $\times$ Aurora (Retail) $\rightarrow$ fit o, suppressed, MiFID reason SHOWN.
4. Mitch's derived goals $\approx$ Growth-dominant ($\sim 70\%+$) from stated-risk; Fotis $\approx$ moderate from revealed.
5. build_today_focus.py exits with ZERO warnings on the shipped payload; untriggered tacticals reported held back.
6. Re-running the build with no payload change moves NO dates (ledger no-op).
7. Editing one idea's thesis bumps only its updatedAt.
8. A chart block without a fetched series is stripped from today_focus.js.
9. Every structures[] string resolves in expressions.js (no dead expression links).
10. Conviction labels: any estimated pillar caps at Medium / "High — data gap".
11. Discarding an idea in Solutions view hides it from the Advisor view; approving restores it; a Solutions "add idea" draft shows a DRAFT pill (no ring), maps to clients via the live engine, and survives reload (localStorage).
12. An Advisor Book tile opens the embedded per-client portfolio report with the shell's theme applied and a working " $\leftarrow$ All clients" return.

What NOT to do. No build tooling, no frameworks, no npm — plain JS files loaded by script tags. No server-side rendering. Never hand-edit today_focus.js. Never fabricate a citation, a consensus number, or an earnings result — the validators treat these as defects, and so does the project's whole reason for existing.

